

# **CSC7970 – NEXT-GENERATION NETWORKING**

## **CONTENT DELIVERY NETWORKS**

**Instructor: Susmit Shannigrahi**  
**[sshannigrahi@tntech.edu](mailto:sshannigrahi@tntech.edu)**

# CDNs

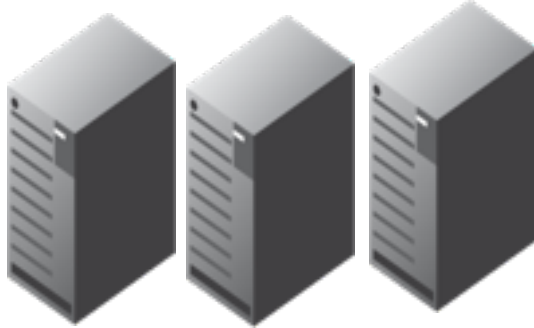
- Content is the primary focus of the Internet
  - What was the Internet designed for?
  - Content distribution requires workarounds
- Main idea
  - Move content near the user using geographically distributed servers
  - Cache popular content in a application overlay
- **Goal: Scalable** content distribution
  - Example: Stranger things

# The CDN model

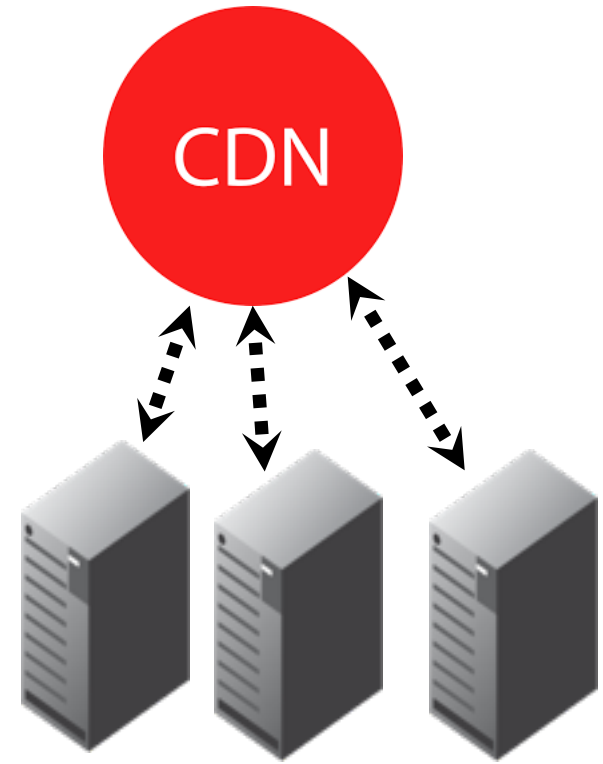
Single server model



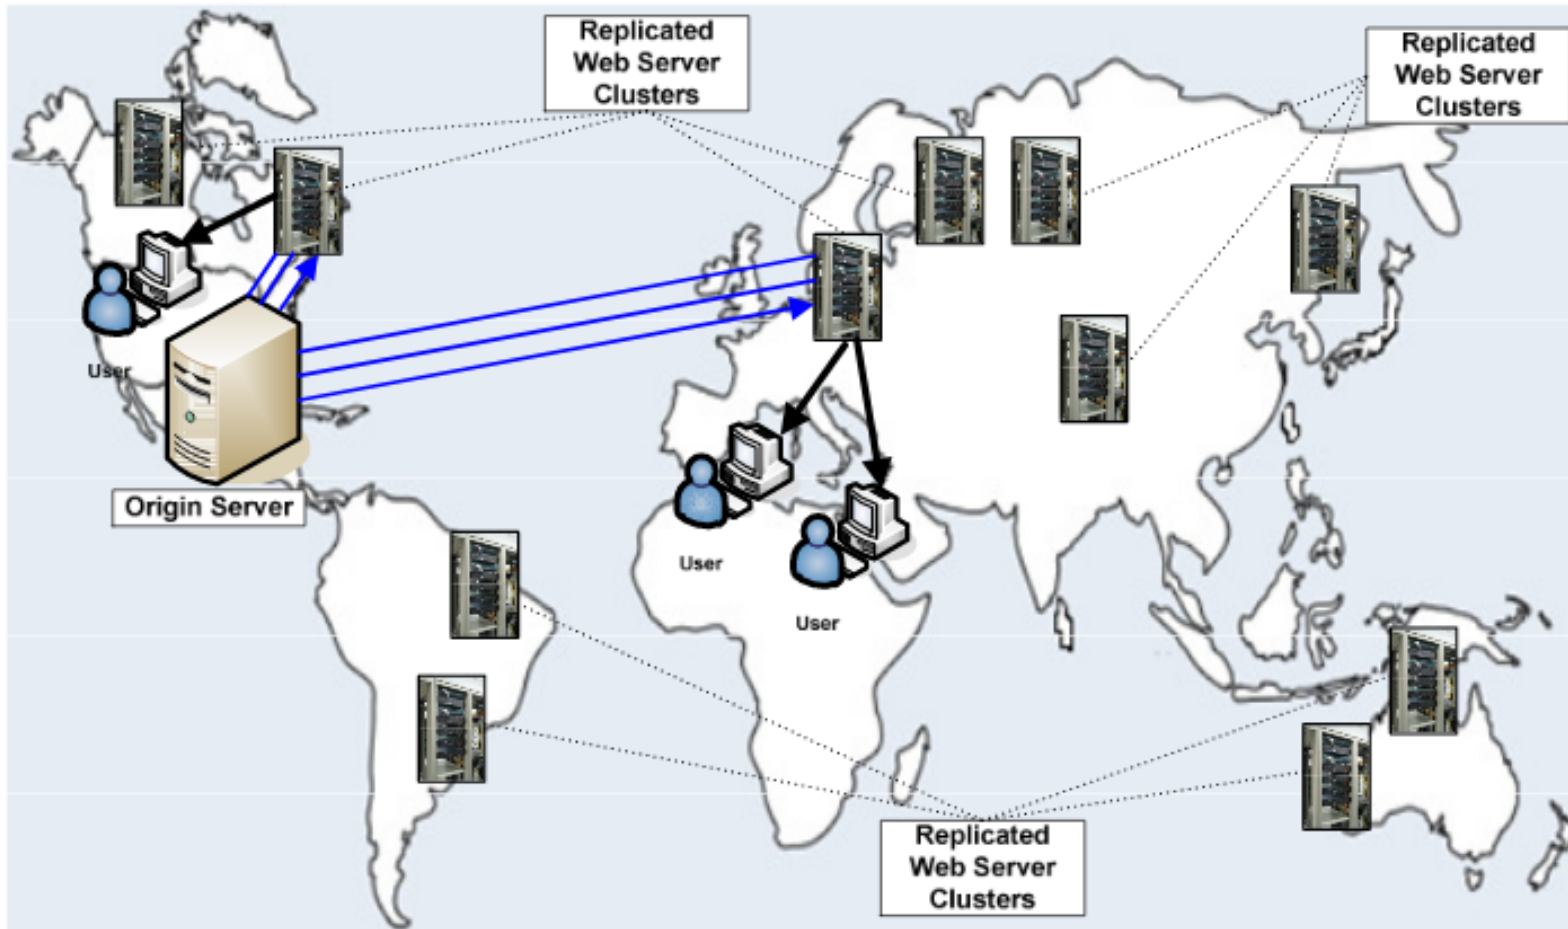
Multi-server model  
(replicated content)



CDN model



# CDN Entities



Origin Server

Replica Servers

Clients

Point of Presence (POP)

Caches

DNS

# Who uses CDNs?

- Everyone
  - But mainly geared towards content distribution
- Why would a content owner use them?
  - TCP/IP Internet architecture does scale for content distribution
  - Load and performance challenges
  - Improved response times

# How important can “latency” be?

- Affects user experience

- Users lose attention if a page takes more than a few hundred ms to load
- 2 Seconds is an eternity

- Affects corporate revenue. Examples:

- *Amazon*: revenue increase of 1% for every 100ms of reduction in page response time
- Google experienced a 20% decrease in ad revenue with a half-second increase in page load time.
- *Shopzilla*: revenue increase of 12% by reducing page response time from 6 seconds to 1.2 seconds

# Who provides CDN services?

- Also everyone:
  - Akamai started it, Amazon, Google, MS, Verizon,
- Big companies have their own CDNs
  - Why?
  - Expensive
  - If a provider fails, their data becomes unavailable
  - Vendor lock-in
  - Costs money to migrate

# Percent of content coming from CDNs

- Any guess/idea?
- *Conservative estimation*: half of the bits of Internet traffic
- *Probably closer to reality*: 80-90% of the bits of Internet traffic



# “Reverse-engineering” CDNs

Spyros\$ dig www.apple.com

; <<>> DiG 9.10.6 <<>> www.apple.com

;; global options: +cmd

;; Got answer:

;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 57705

;; flags: qr rd ra; QUERY: 1, ANSWER: 4, AUTHORITY: 0, ADDITIONAL: 0

;; QUESTION SECTION:

;www.apple.com. IN A

;; ANSWER SECTION:

www.apple.com. 57 IN CNAME www.apple.com.edgekey.net.

www.apple.com.edgekey.net. 16510 IN CNAME www.apple.com.edgekey.net.globalredir.akadns.net.

www.apple.com.edgekey.net.globalredir.akadns.net. 1975 IN CNAME e6858.dsce9.akamaiedge.net.

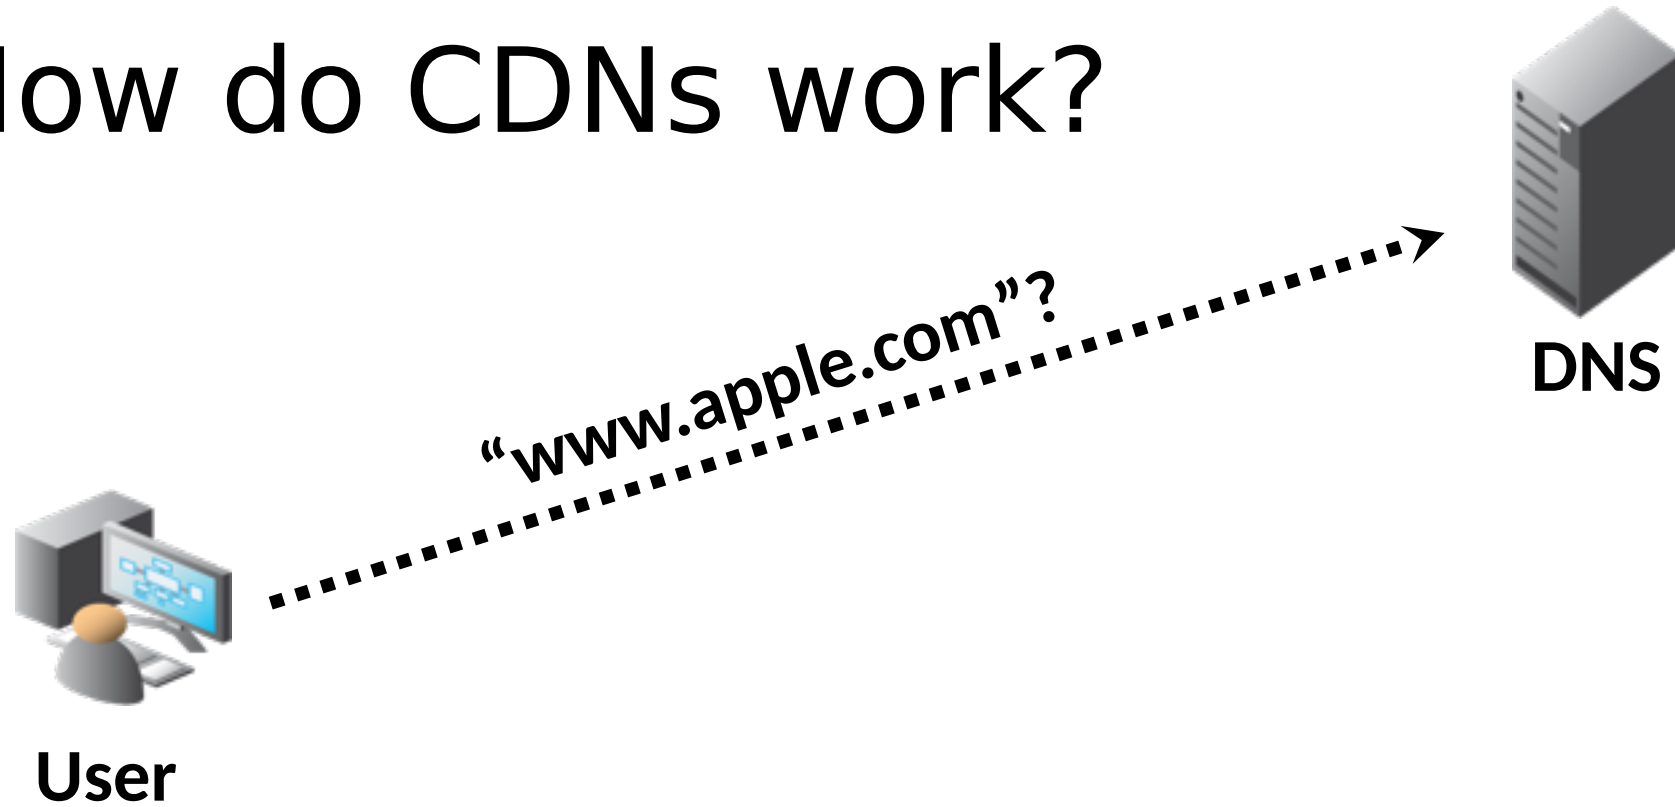
e6858.dsce9.akamaiedge.net. 14 IN A 23.64.139.135

# How do CDNs work?

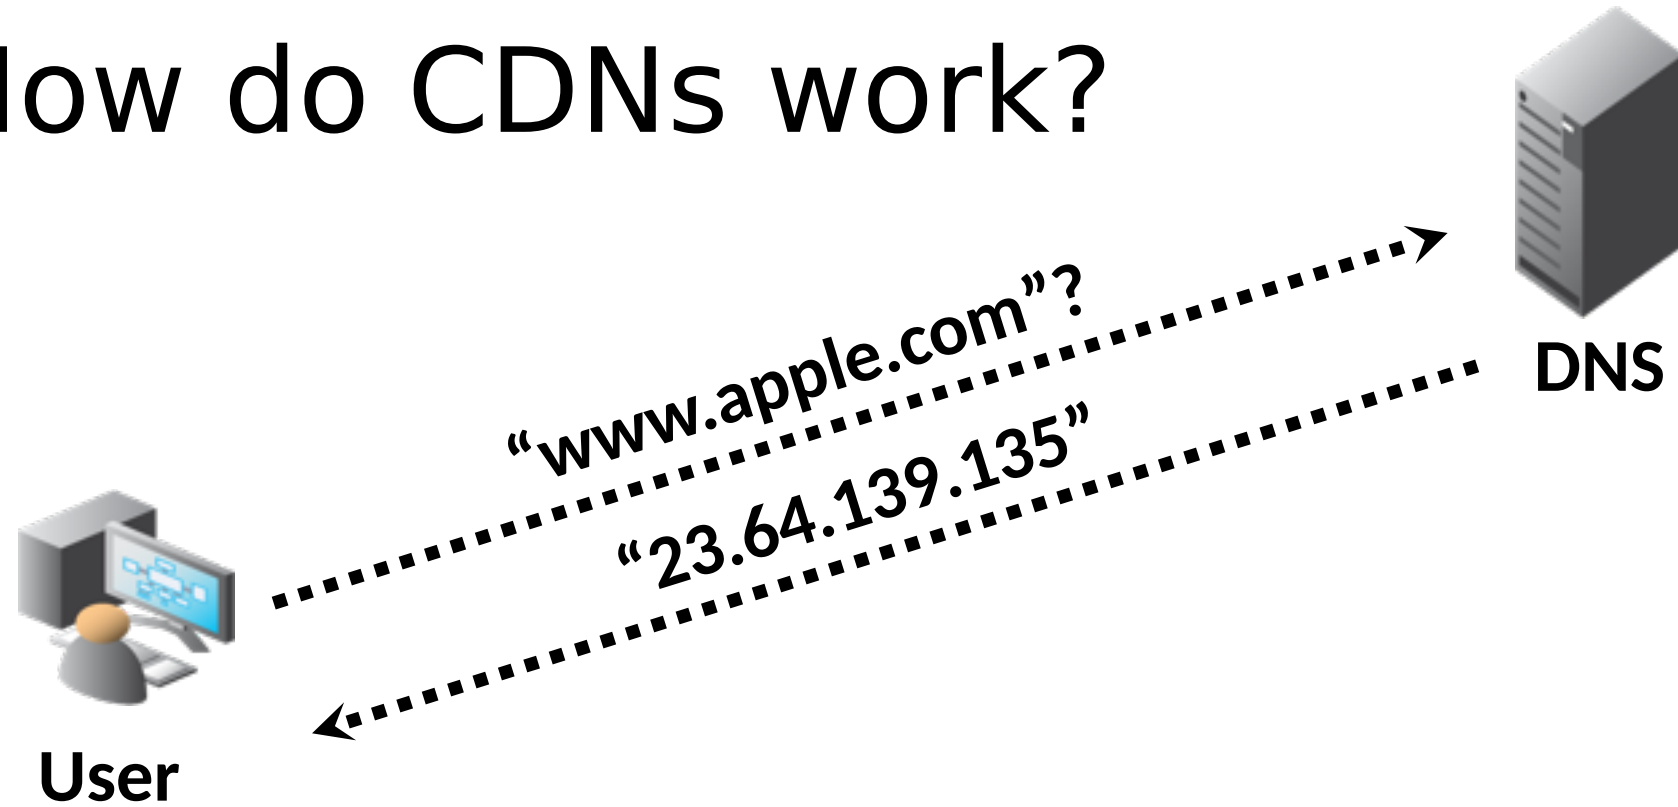


**User**

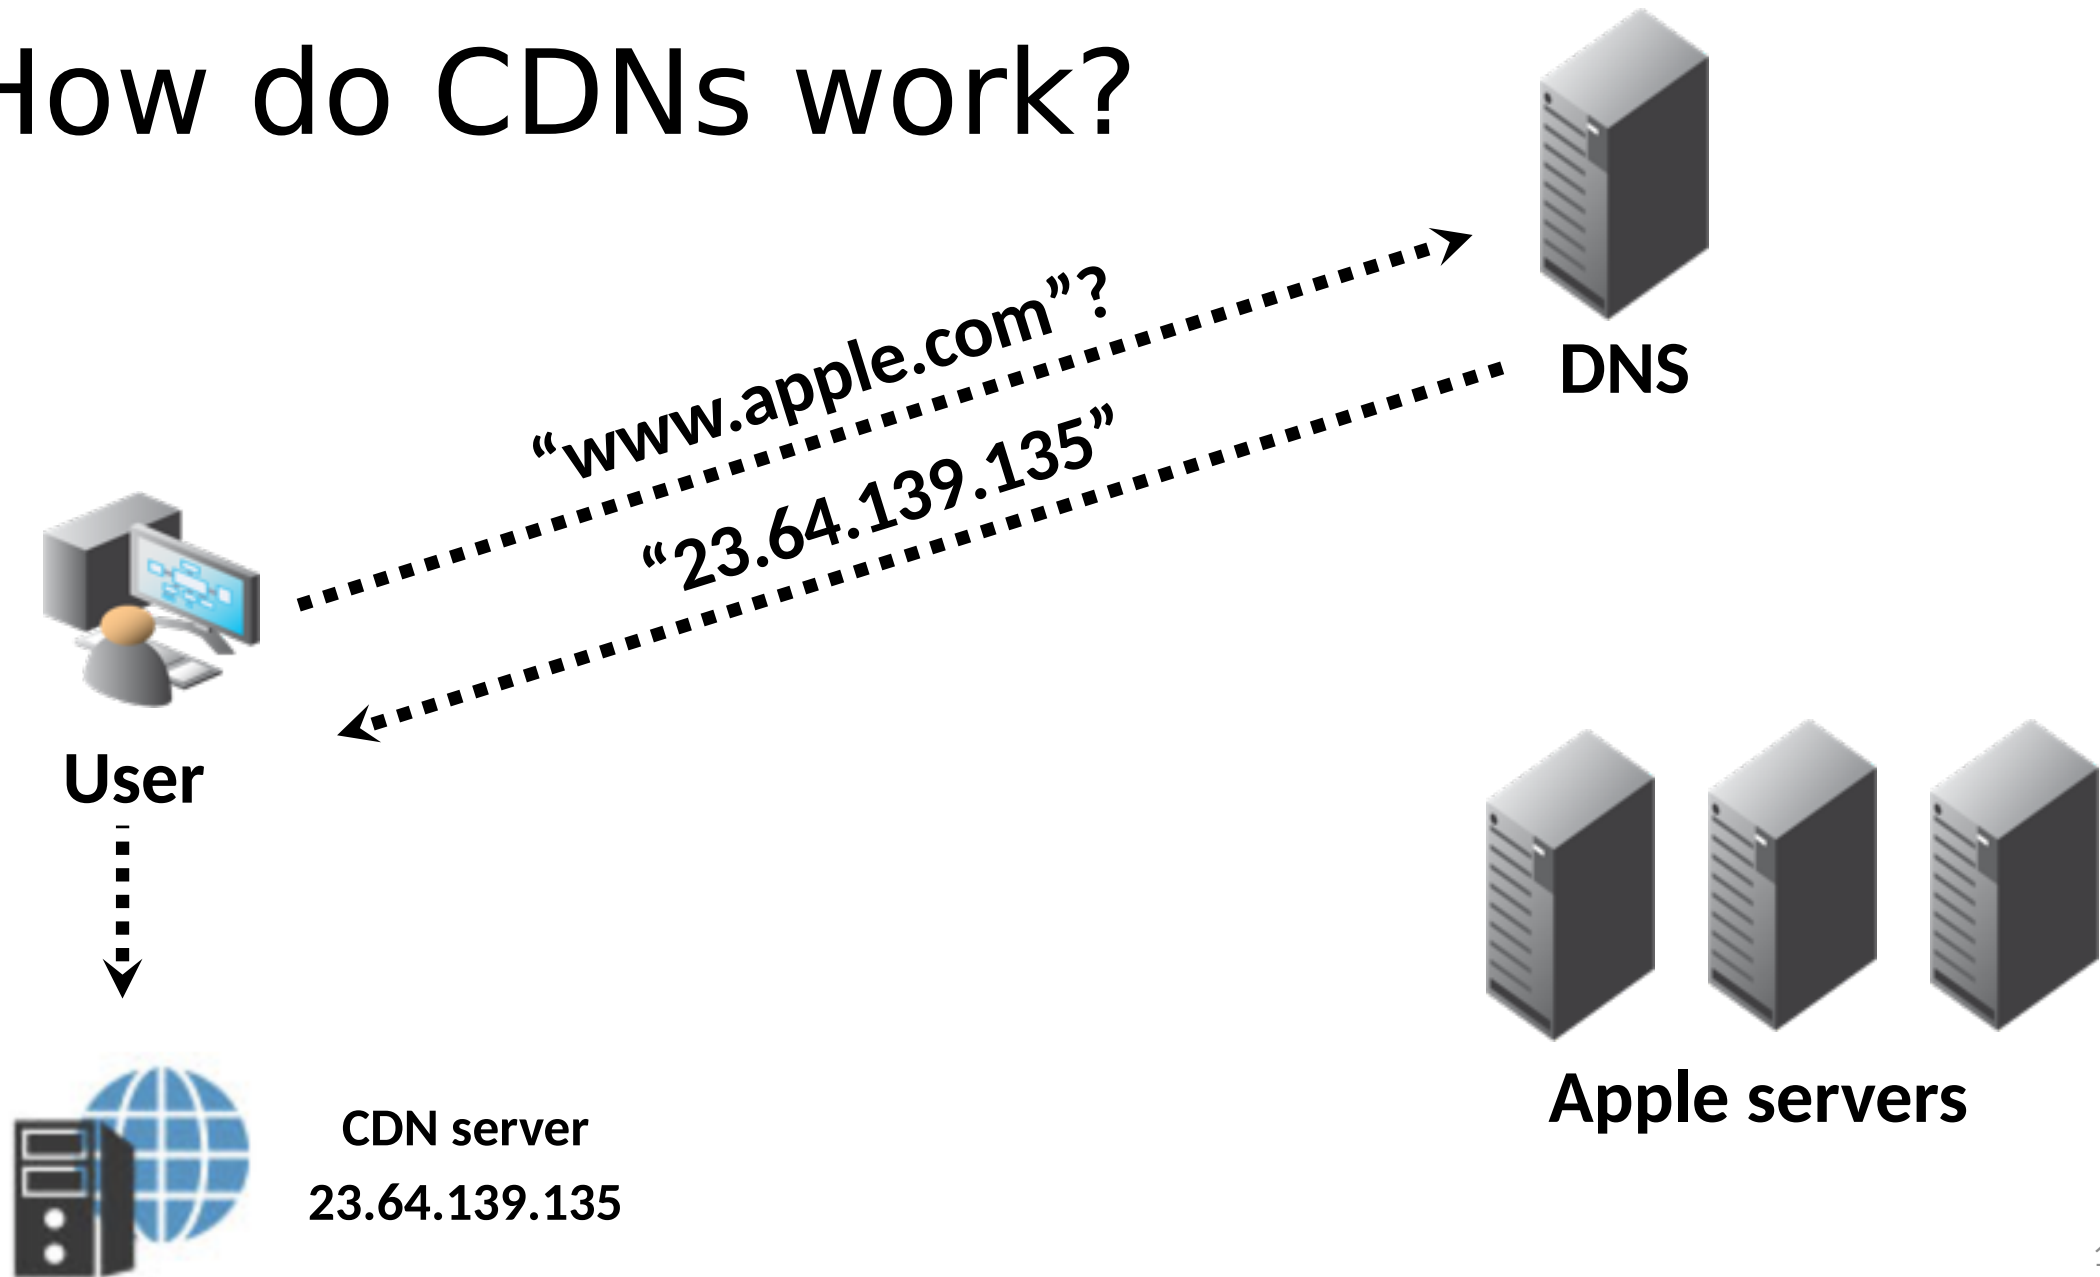
# How do CDNs work?



# How do CDNs work?



# How do CDNs work?



# Discovering a CDN server

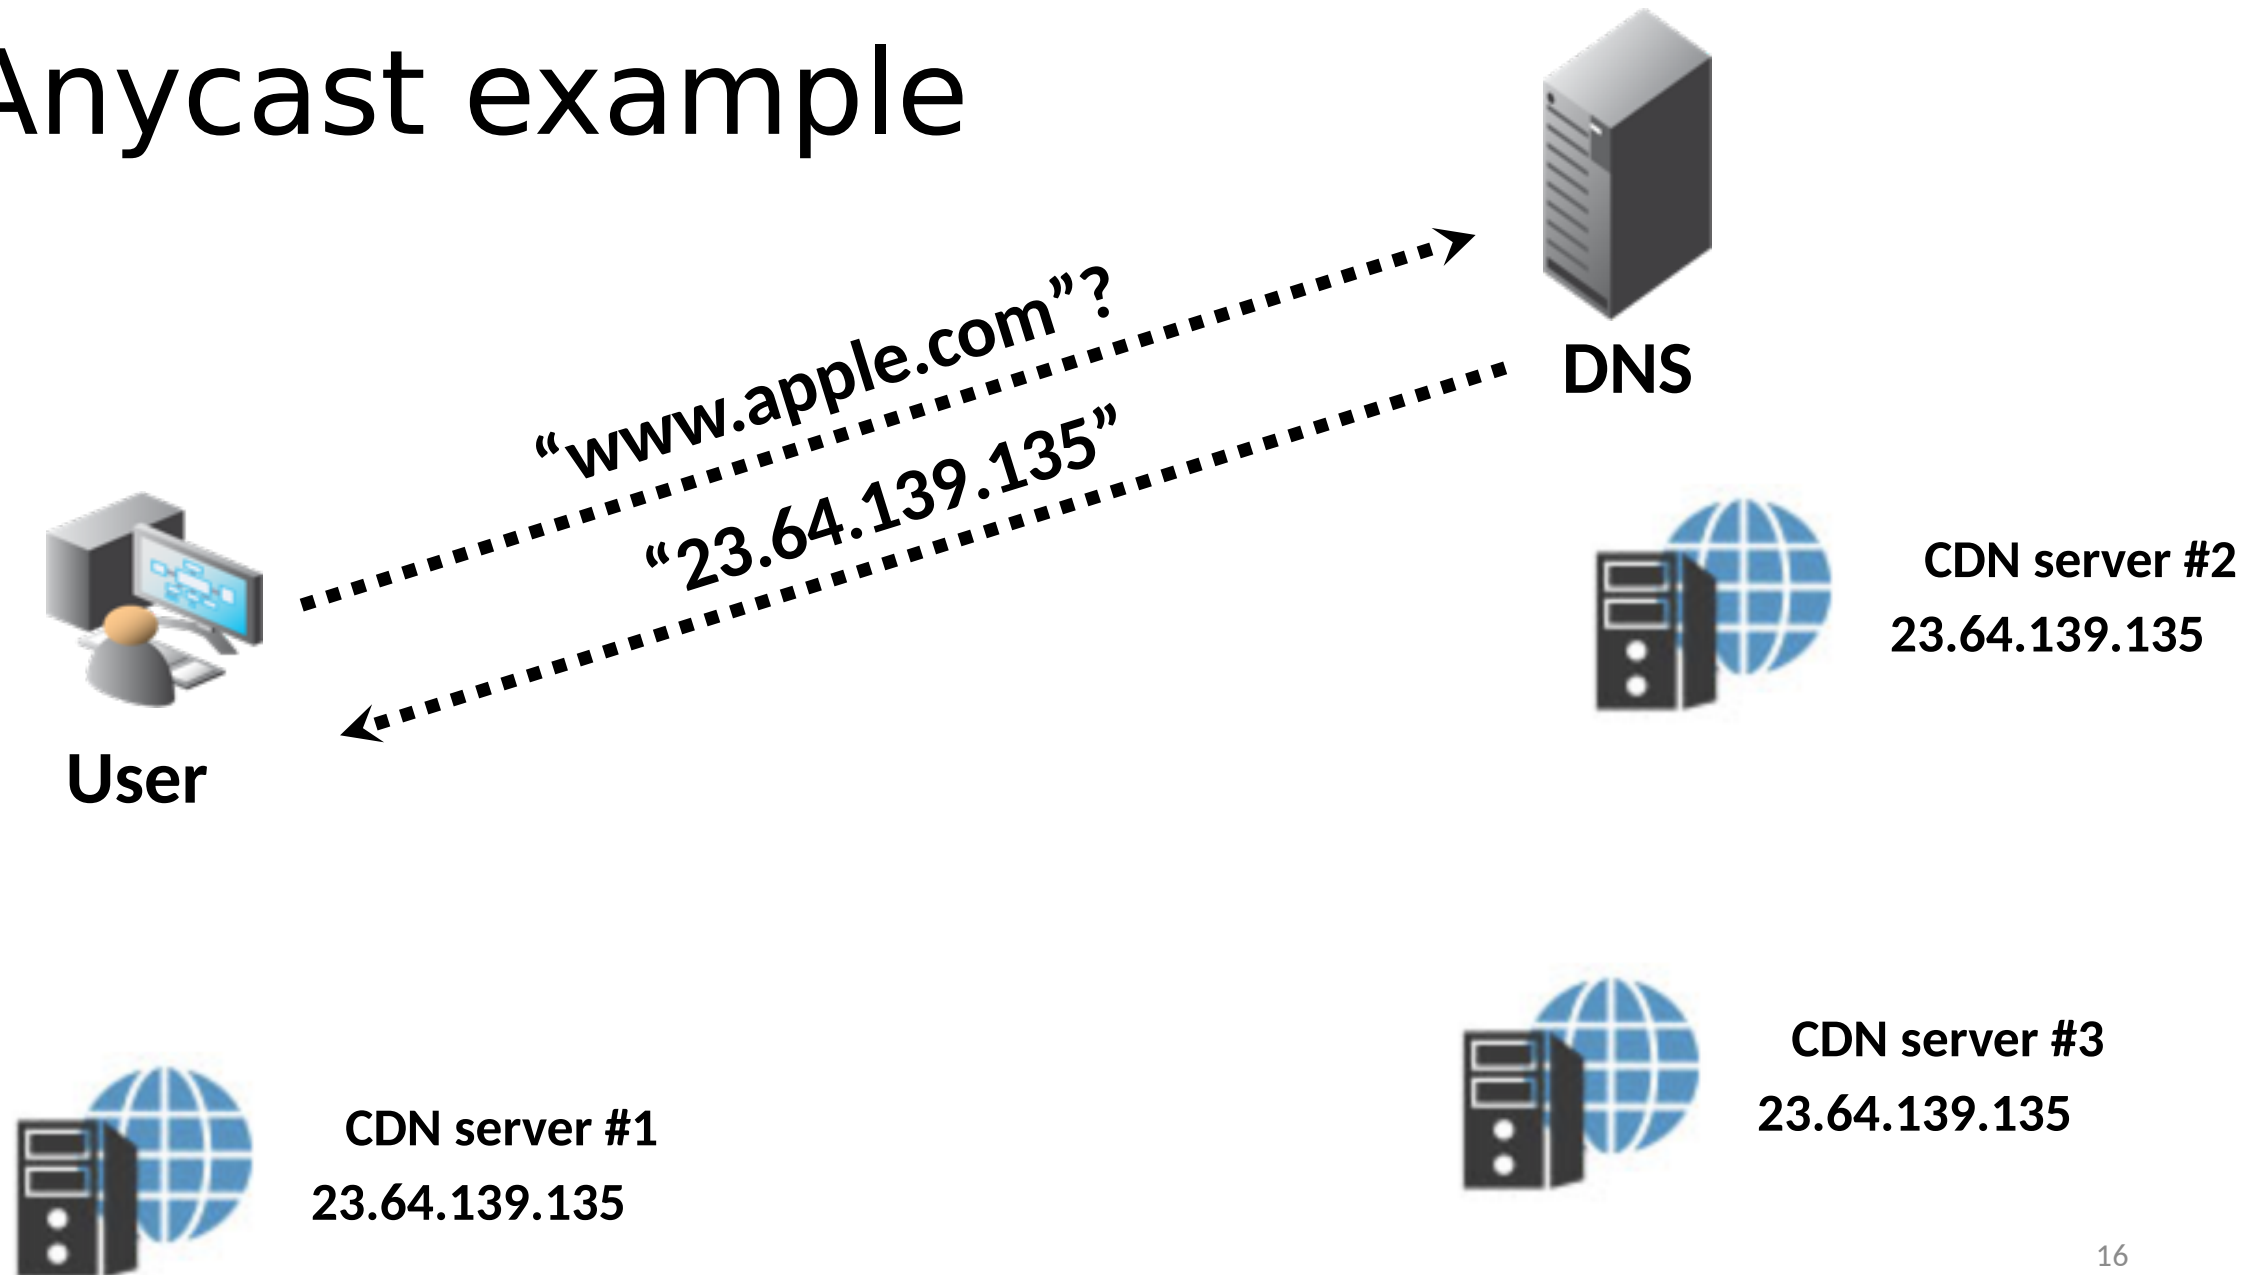
- DNS redirection
  - DNS returns IP address of CDN server(s) to clients
- Advantages
  - Any ideas?
  - DNS already has scalable infrastructure
  - URLs need no change
- Limitations
  - Who controls the user content?
  - DNS typically returns a single IP address
  - This IP address might be unicast

# How to get around the DNS limitations?

- IP anycast

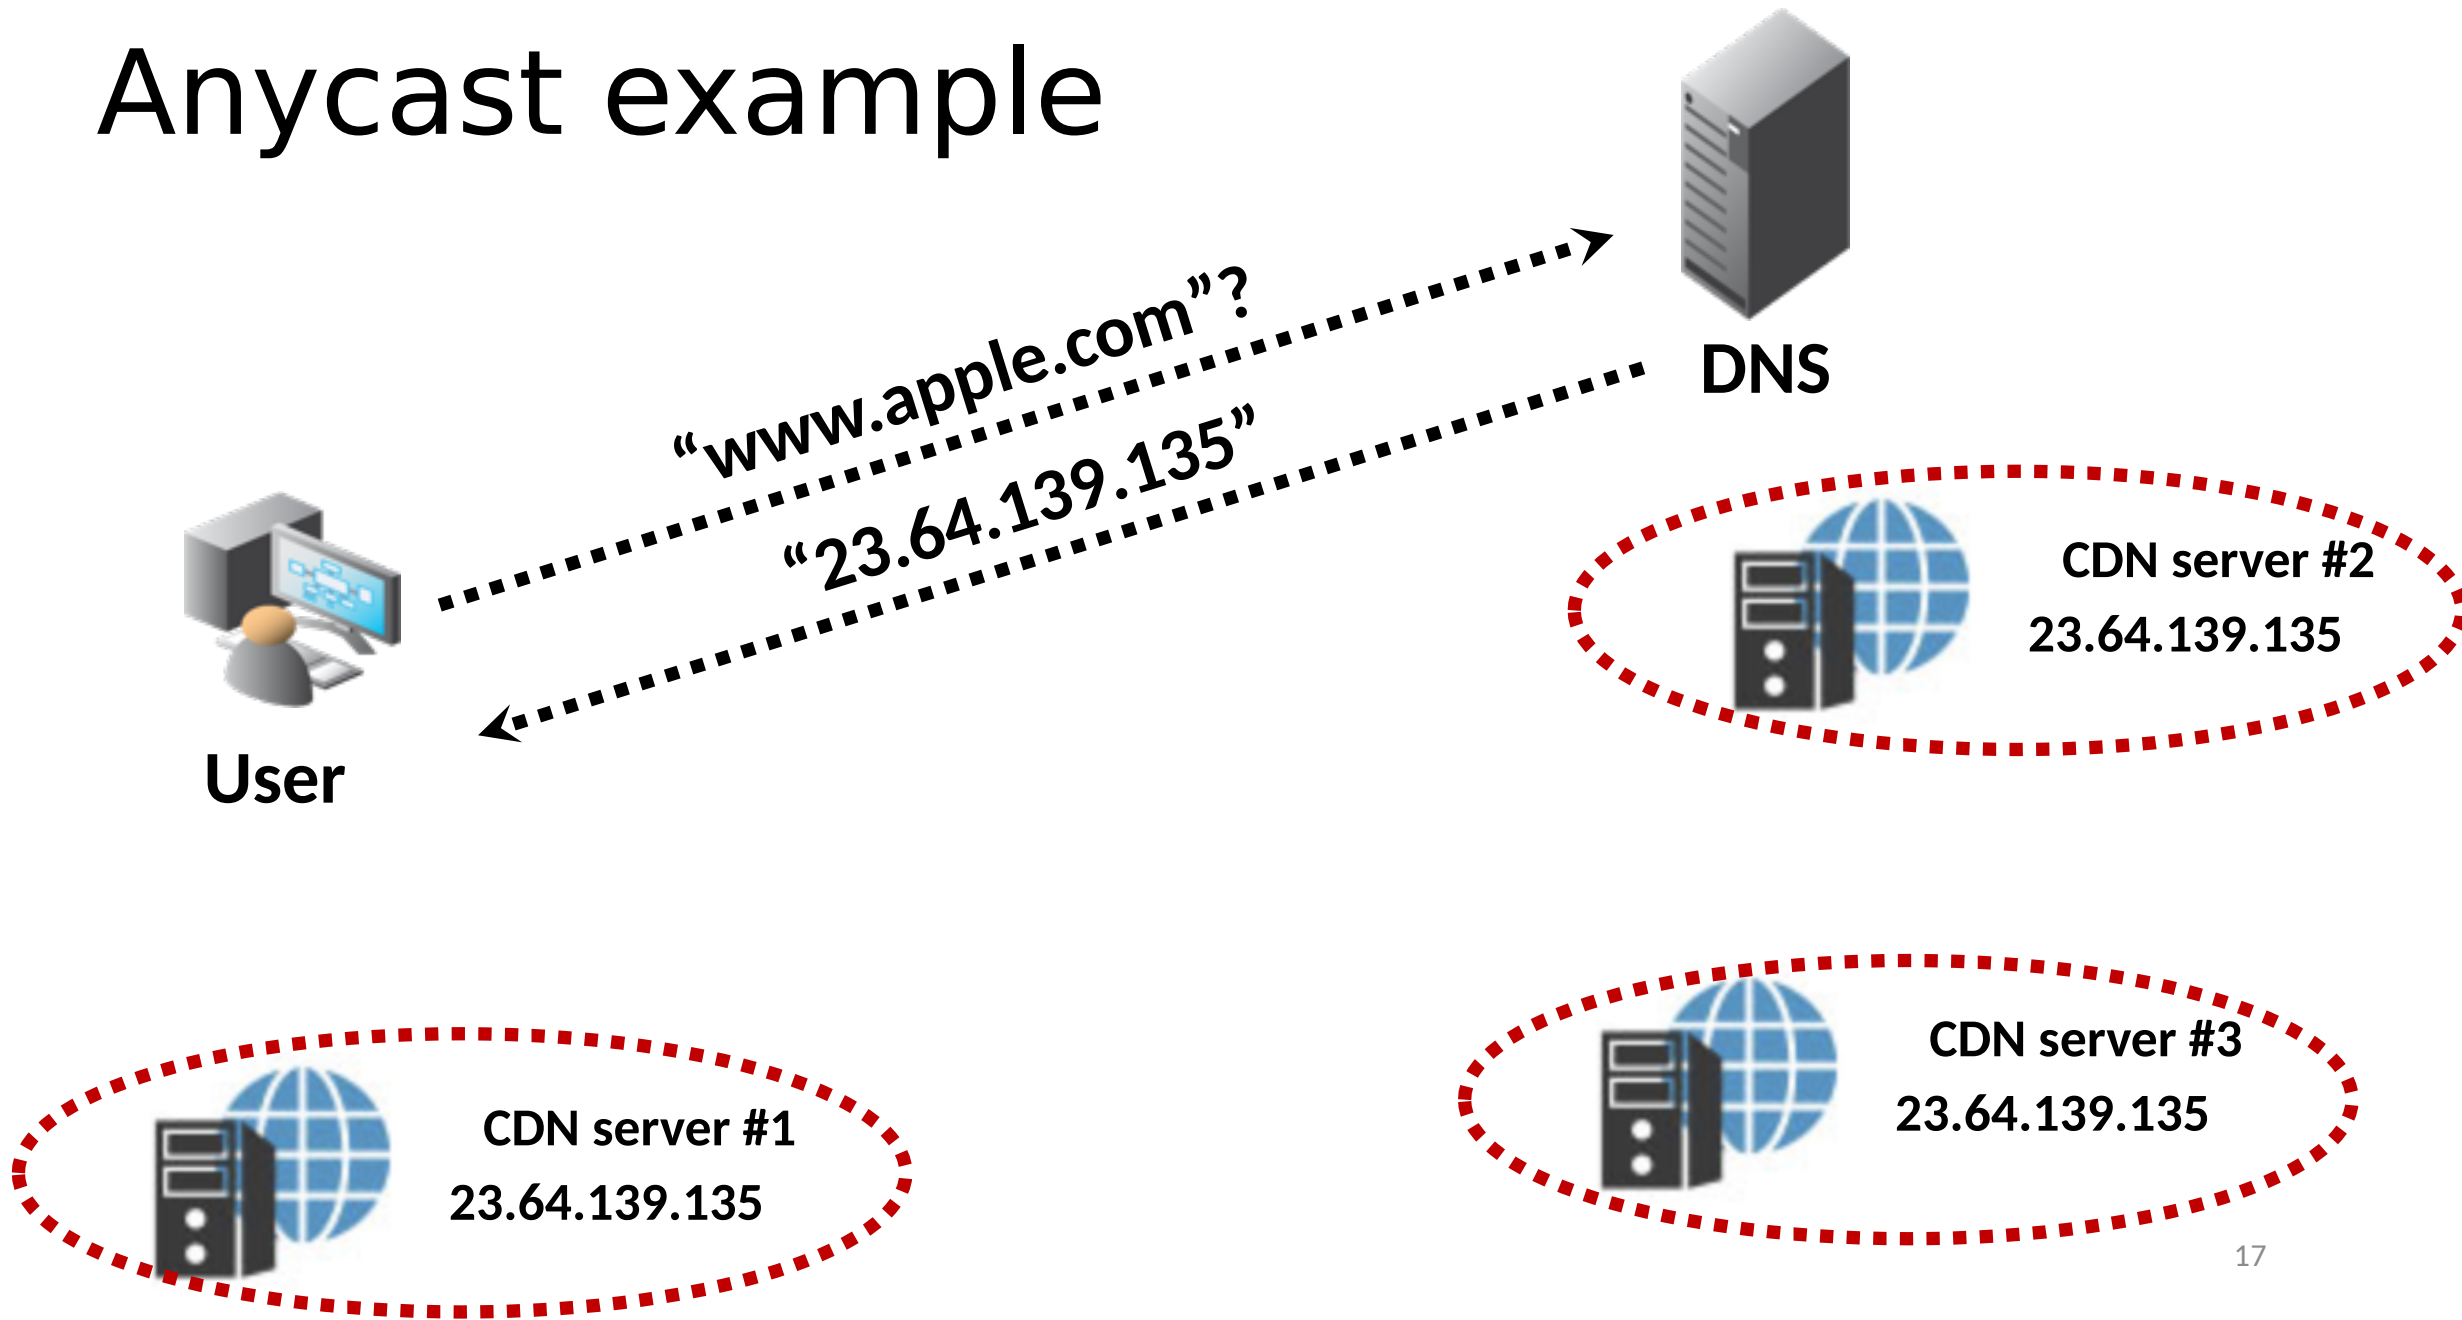
- Deploy multiple CDN servers with same IP address
- All these servers have the requested content
- Content requests reach the closest CDN server to the user

# Anycast example

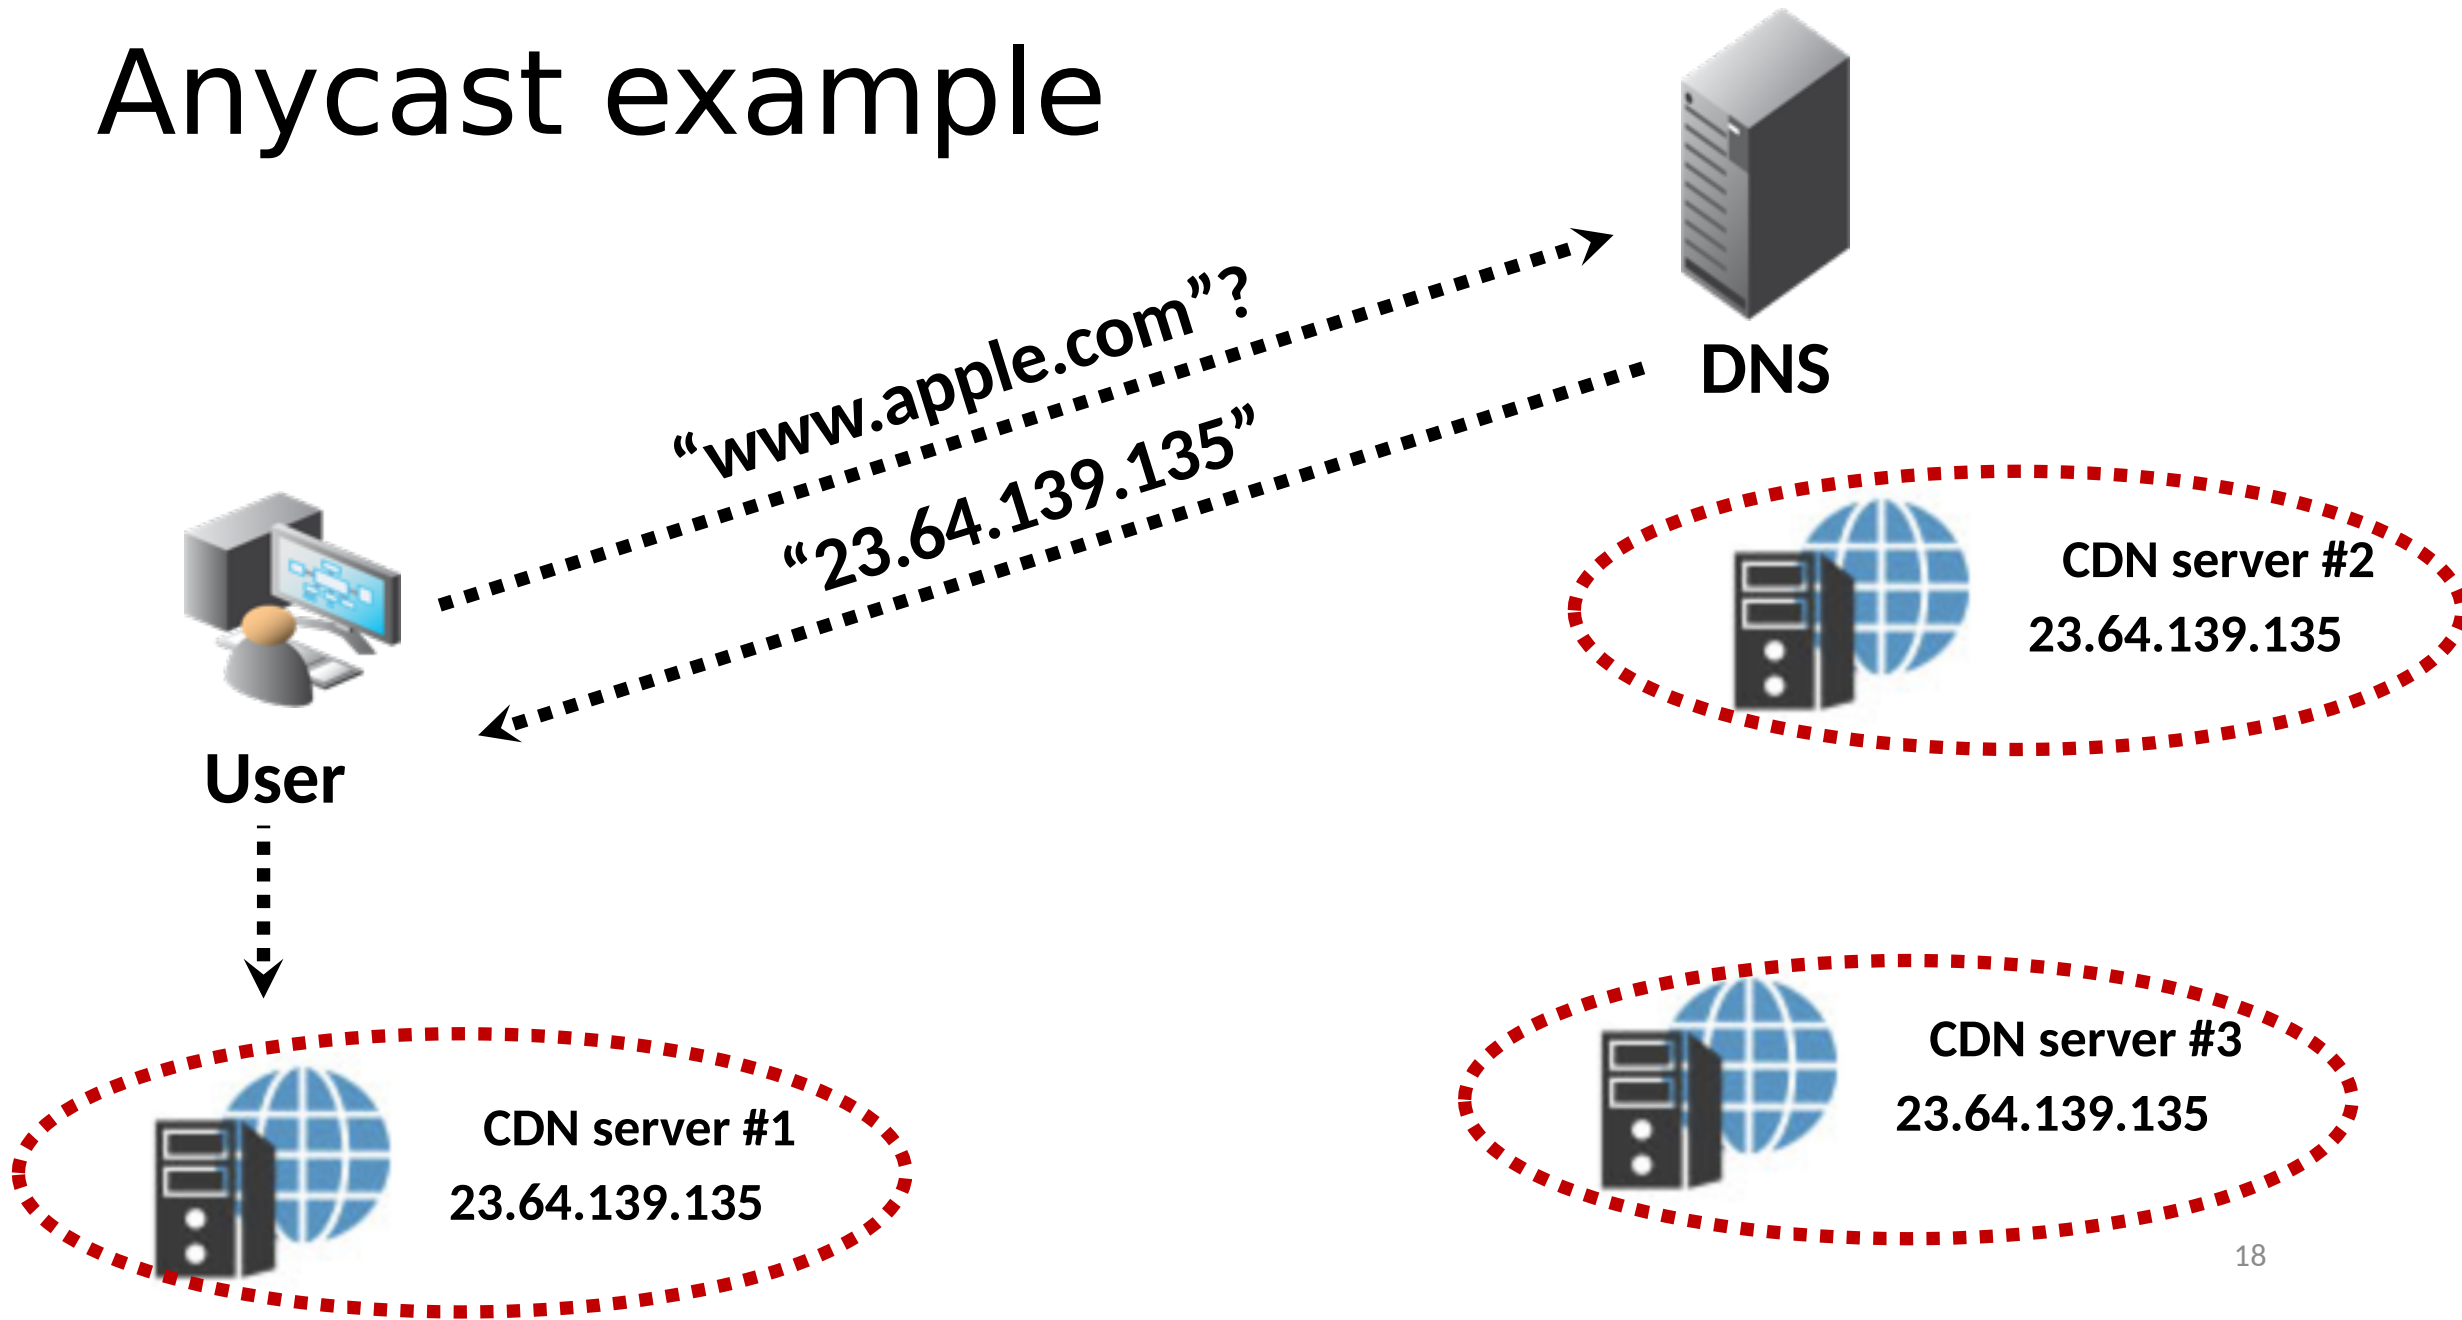




# Anycast example



# Anycast example



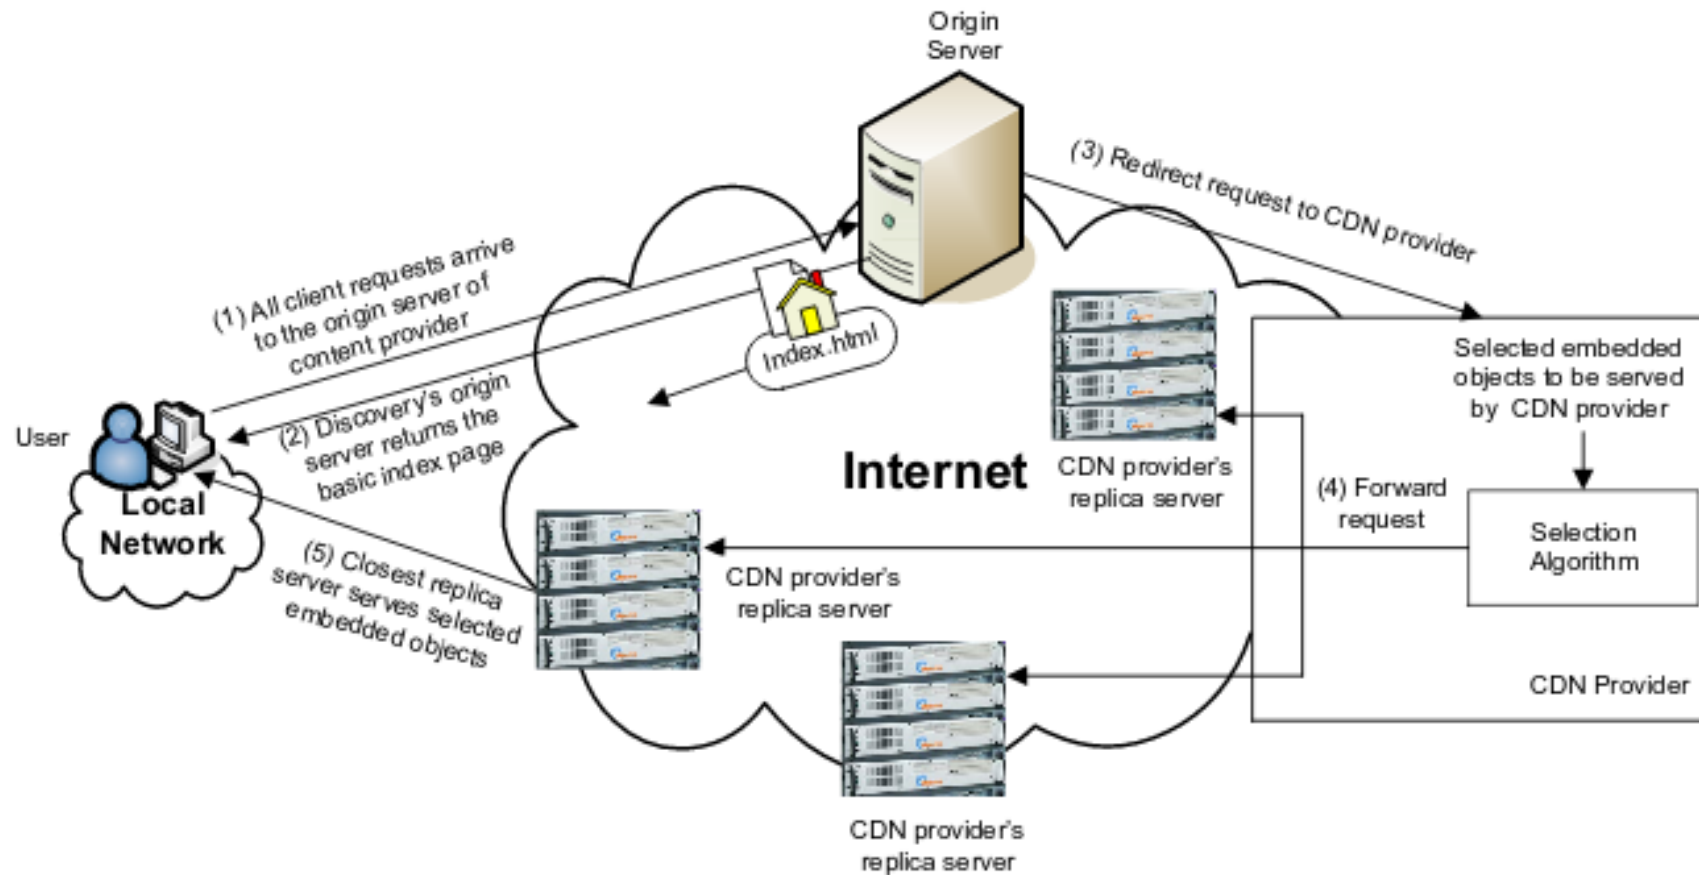
# Further optimizing on performance

- Multiple CDN servers in same data-center having the requested content
  - How to pick the "best" one?
- It depends on how we define "the best one"
  - Load, network conditions, service level agreements
  - Typically solved through load balancers, network monitoring, and configuration

# Request Routing

- Multiple CDN servers in same data-center having the requested content
  - How to pick the "best" one? ← Request routing
- It depends on how we define "the best one"
  - Load, network conditions, service level agreements
  - Typically solved through load balancers, network monitoring, and configuration

# Request Routing



# How CDNs handle dynamic content?

- What to do with dynamic content?
  - Examples: webpages with ads or websites that require passwords
  - Break dynamic pages into parts that can be cached and parts that cannot
  - CDN server fetches non-cacheable parts (e.g., response to a password) from the original content server and cacheable parts from its own cache
  - CDN server puts non-cacheable and cacheable parts together, assembles the website, and sends it back to the user
- How effective could that be?
  - According to Akamai, this reduces bandwidth requirements for dynamic content by 95-99%

# Security/Privacy

- Impact on security/privacy?
  - Users and content providers need to trust CDNs!
- Who has control over user content?
  - For sure, not the user!
- What happens if the CDN gets compromised?
  - Content poisoning!

# Encrypted content

- Why not using end-to-end encryption to solve all the security/privacy problems?
  - CDNs need to have access to unencrypted content to optimize content delivery!
- *User think they establish an encrypted connection to the content server!*
  - Typically NOT true!
  - CDNs usually have access to encryption/decryption keys!
  - CDNs fetch content from content servers --> store content (typically unencrypted) --> encrypt requested content and provide it to users



# Mitigate DDoS attacks through CDNs

- Lately, DDoS attacks are against applications/services (e.g., Facebook) in addition to network services (e.g., DNS)
  - Attack a large distributed network of caches harder than attacking a few content servers directly!
  - Attacks against CDNs (big guys) instead of content owners (smaller guys)
  - CDNs monitor traffic -> detect anomalies and attacks
  - Drop or reroute malicious traffic almost in real-time

# Conclusion

- CDNs are ubiquitous now
- CDNs are still under active development
- Many free CDNs you can use today
  - Cloudflare, netli and many others

# Akamai Paper

- Main takeaways

- CDNs started for solving the flash crowd problem
- Multi-homing, clusters, mirroring do not scale
- Define network functions – nearest, available, and likely
- DNS + Defined Network Functions for choosing servers
- Dynamically assemble content for delivery
- For streaming content – get data from origin server and send to the edge, users get content from the edge

- Massively replicate content and services

- Versioned objects for content delivery

- Close communication with content providers (trust)

# Netflix Paper

- Main takeaways

- How to measure distributed infrastructure (again)
- Netflix owns its CDN – cost
- Distribute a plugin and convince people to use it, look at the requests and traffic
- Server names are location specific and has a cache numner
  - ipv6\_1-lagg0-c002.1.lhr005.bt.isp.nflxvideo.net
  - Well defined and hierarchical
- Massive scale
  - ~4300 servers in the US alone! (Mostly at IXPs)
  - No deployment in major ISPs – they want Netflix to pay for traffic
  - In Brazil – the opposite – mostly at ISPs – why?
- What about complexity?