

# **CSC4200/5200 – COMPUTER NETWORKING**

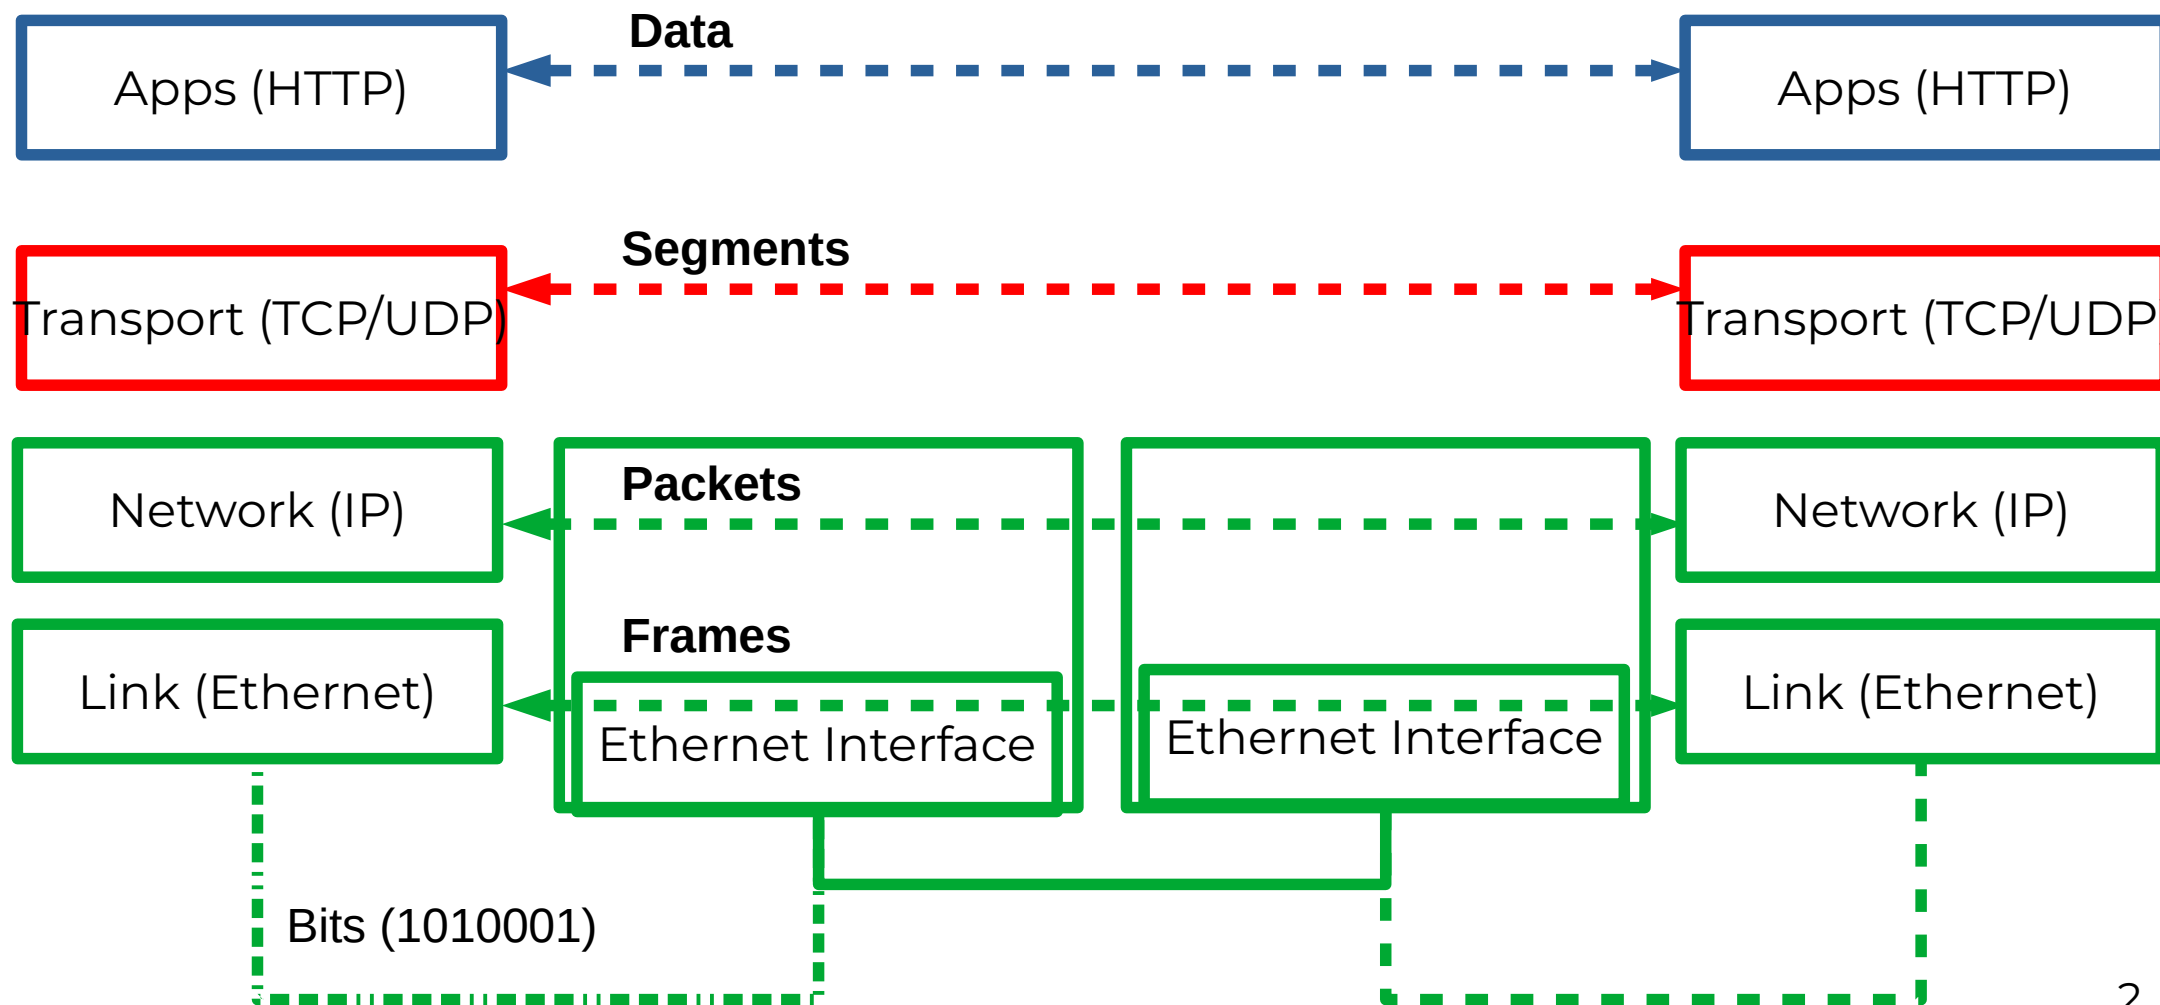
**Instructor: Susmit Shannigrahi**

**CONGESTION CONTROL**

**sshannigrahi@tnitech.edu**


**GTA: dereddick42@students.tnitech.edu**





# TCP flow control

- receiver “advertises” free buffer space in the header
- sender limits amount of unacked (“in-flight”) data to receiver’s **rwnd** value
- guarantees receive buffer will not overflow



# Congestion Control




# Principles of congestion control

## *congestion:*

- informally: “too many sources sending too much data too fast for *network* to handle”
- different from flow control!
- manifestations:
  - lost packets (buffer overflow at routers)
  - long delays (queueing in router buffers)
- a top-10 problem!


# Congestion: scenario 1

- [redacted]
- three senders, two receivers
- one router, infinite buffers
- output link capacity:  $R$
- The router can only transmit one ... and either buffer or drop the other
- If many packets arrive,
- Buffer overflow



# Causes/costs of congestion: scenario 2

- one router, *finite* buffers
- sender retransmission of timed-out packet
  - application-layer input = application-layer output:  $\lambda_{in} = \lambda_{out}$
  - transport-layer input includes *retransmissions*:  $\lambda_{in} \neq \lambda_{out}$



# Metrics: Throughput vs Delay

High throughput –


- Throughput: measured performance of a system –E.g., number of bits/second of data that get through
- Low delay –
- Delay: time required to deliver a packet or message –E.g., number of ms to deliver a packet •
- These two metrics are sometimes at odds –
  - More packets = more queuing



# Issues in Resource Allocation

- Evaluation Criteria
  - Effective Resource Allocation

*power of the network.*  
 $\text{Power} = \text{Throughput}/\text{Delay}$



Ratio of throughput to delay as a function of load

# Issues in Resource Allocation

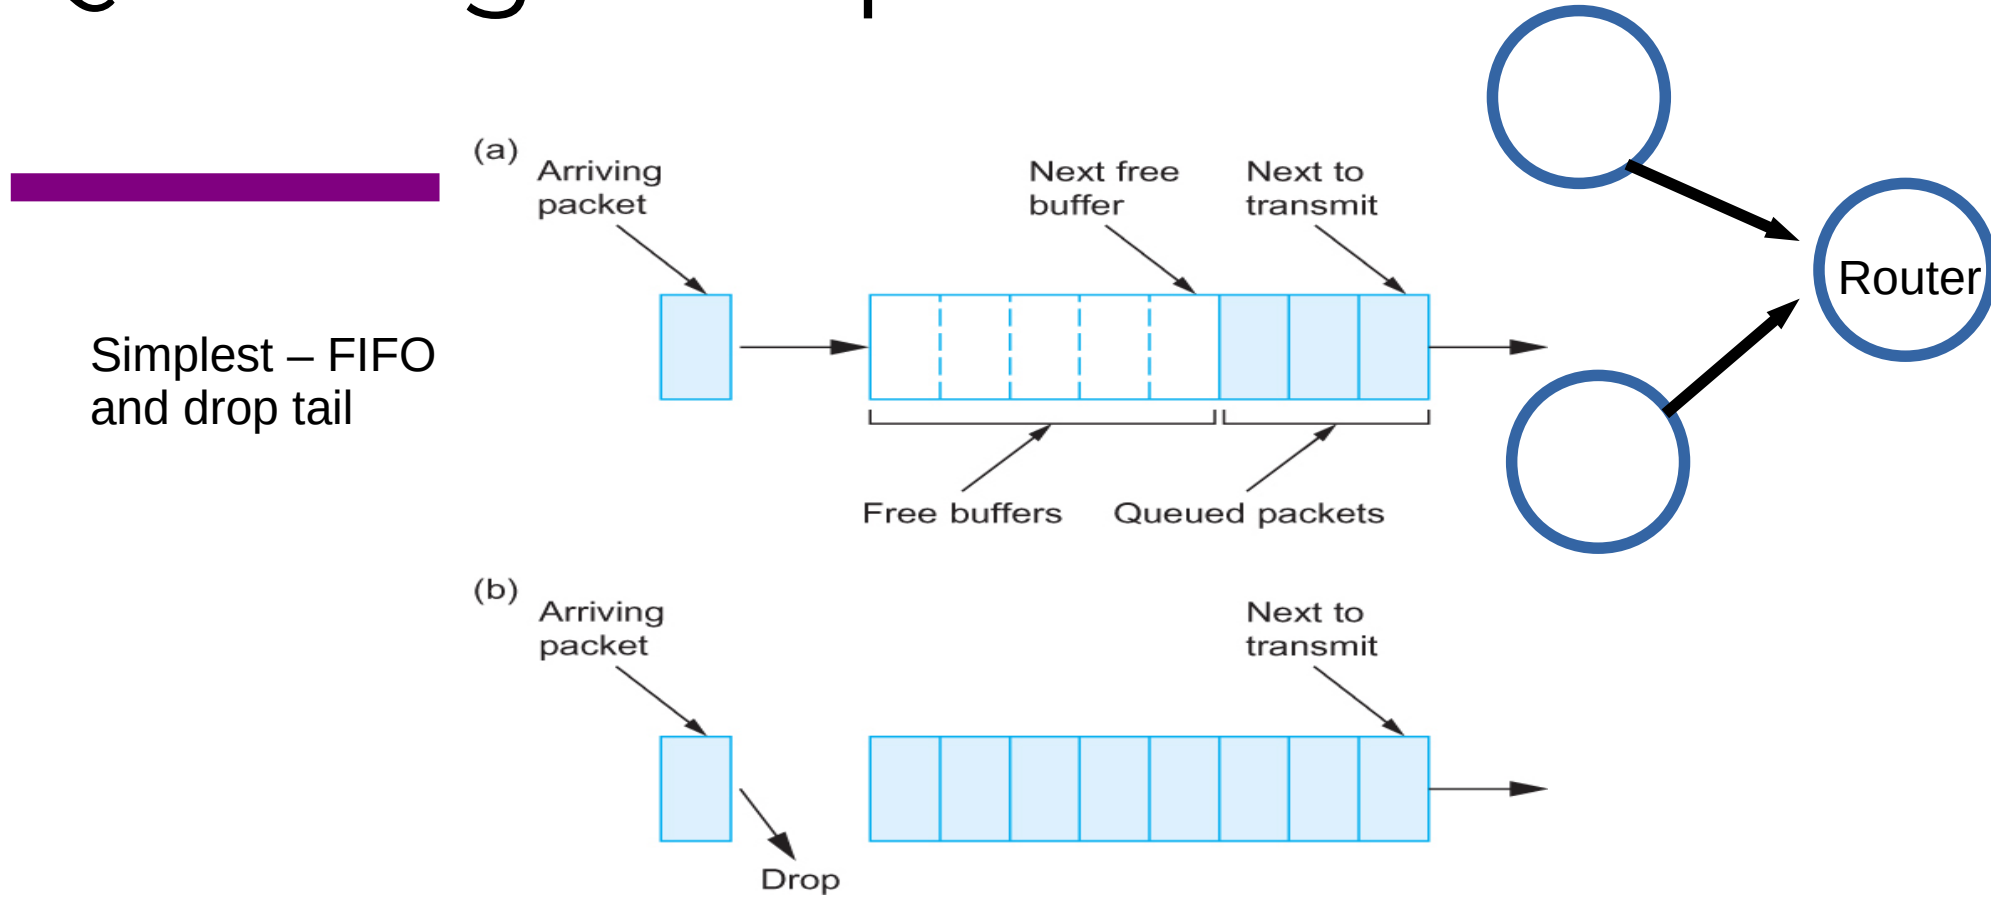
---

- Evaluation Criteria
  - Fair Resource Allocation
    - The effective utilization of network resources is not the only criterion for judging a resource allocation scheme.
    - We want to be “fair”
    - Equal share of bandwidth

But, what if the flows traverse different paths?

Open problem, often determined by economics

# Queuing Disciplines



(a) FIFO queuing; (b) tail drop at a FIFO queue.

What are the problems?

# Defining Fairness: Flows

---

“fair” to whom? – Should be Fair to a Flow

What is a flow?

Combination of <Src IP, Src Port, Dst IP, Dst Port>


# Fair Queuing

---

- Fair Queuing
  - FIFO does not discriminate between different traffic sources, or
  - it does not separate packets according to the flow to which they belong.
  - Fair queuing (FQ) maintains a separate queue for each flow

# Queuing Disciplines

- Fair Queuing



Round-robin service of four flows at a router

# Min Max Fair queuing

---

- Assume  **$n$**  clients
- Channel capacity  **$C$**
- Give  **$C/n$**  to each client
  - If  $C_1$  does not want  $C/n$
  - Divide the excess capacity equally among others
  - So everyone else gets  $C/n + (C/n - C_1)/(n-1)$
  - Repeat for  $C_2$  and others

# TCP Congestion Control

---

- Each source determines available capacity
- Max many packets is allowed to have in transit - window
- Congestion window = # of unacked bytes
- $\text{MaxSendWindow} = \min(\text{congestion window}, \text{receiver window})$
- How do you change congestion window?
  - Decrease on losing a packet (back off)
  - Increase on successful send



# How much to increase and decrease?


---

- Additive Increase, Multiplicative Decrease (AIMD)

# How much to increase and decrease?

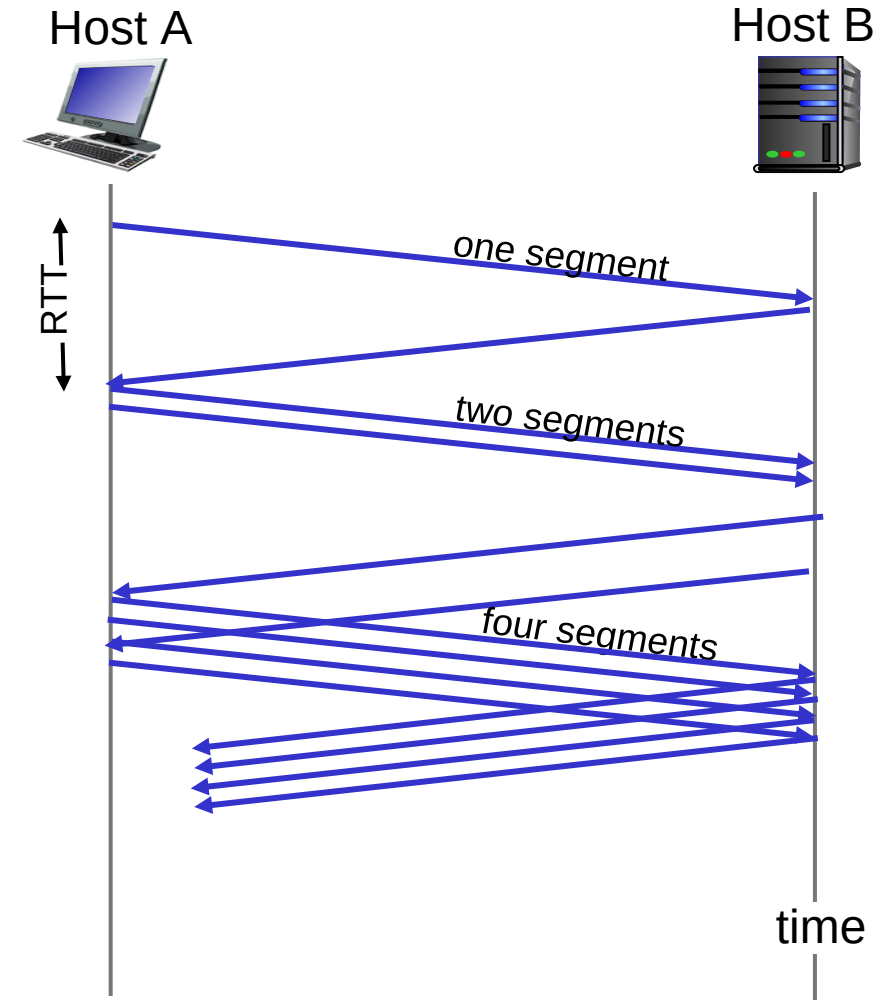
- ❖ *approach*: sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
  - *additive increase*: increase **cwnd** by 1 MSS every RTT until loss detected
  - *multiplicative decrease*: cut **cwnd** in half after loss

AIMD saw tooth behavior: probing for bandwidth



# TCP Slow Start


- when connection begins, increase rate exponentially until first loss event:
  - initially **cwnd** = 1 MSS
  - double **cwnd** every RTT
  - done by incrementing **cwnd** for every ACK received
- summary: initial rate is slow but ramps up exponentially fast



# TCP Slow Start

Why not start with a large window?

Why not increase one by one?



# TCP: detecting, reacting to loss

---


- loss indicated by timeout:
  - **cwnd** set to 1 MSS;
  - window then grows exponentially (as in slow start) to threshold, then grows linearly
- loss indicated by 3 duplicate ACKs: TCP RENO
  - dup ACKs indicate network capable of delivering some segments
  - **cwnd** is cut in half window then grows linearly
- TCP Tahoe always sets **cwnd** to 1 (timeout or 3 duplicate acks)

# TCP:Two types of loss

---

- Triple duplicate ack
  - Do a multiplicative decrease, keep going
- Timeout
  - Reset CWND to 1
  - Take advantage of

# TCP Slow Start and congestion avoidance



How to set ssthresh?

Initially – Randomly high

Later – adjusted as congestion happens


# TCP Congestion Summary

**CWND < Threshold** → Slow Start, Exponential increase

**CWND > Threshold** → Congestion Avoidance, Linear increase

**Triple Duplicate ACK** → Threshold = CWND/2, CWND = CWND/2

**Timeout** → Threshold = CWND/2, CWDN = 1 (or 3)





# TCP Throughput

---

TCP average throughput as a function of window size and RTT?  
Ignore slow start, assume long TCP flow

Let  $W$  be the window size

Throughput =  $W/\text{RTT}$

After loss, throughput =  $W/2 \cdot \text{RTT}$

Average throughput =  $0.75W/\text{RTT}$

# Problems with Fast Links

---

Consider the high speed link:

9000 byte segments

100ms RTT

100Gbps/second throughput

Throughput =  $0.75W/\text{RTT}$

So, WindowSize (w) = Throughput \* RTT / 0.75

$W = 1,481,481,444$  segments

# Problems with Fast Links



TCP assumes all losses are due to congestion

Throughput =  $(1.22 * MSS) * (RTT / \sqrt{Loss})$

What is the loss rate to maximize 100Gbps pipe with 9000 bytes segments and 100ms RTT? Hint – must be very very low

[https://www.switch.ch/network/tools/tcp\\_throughput/](https://www.switch.ch/network/tools/tcp_throughput/)